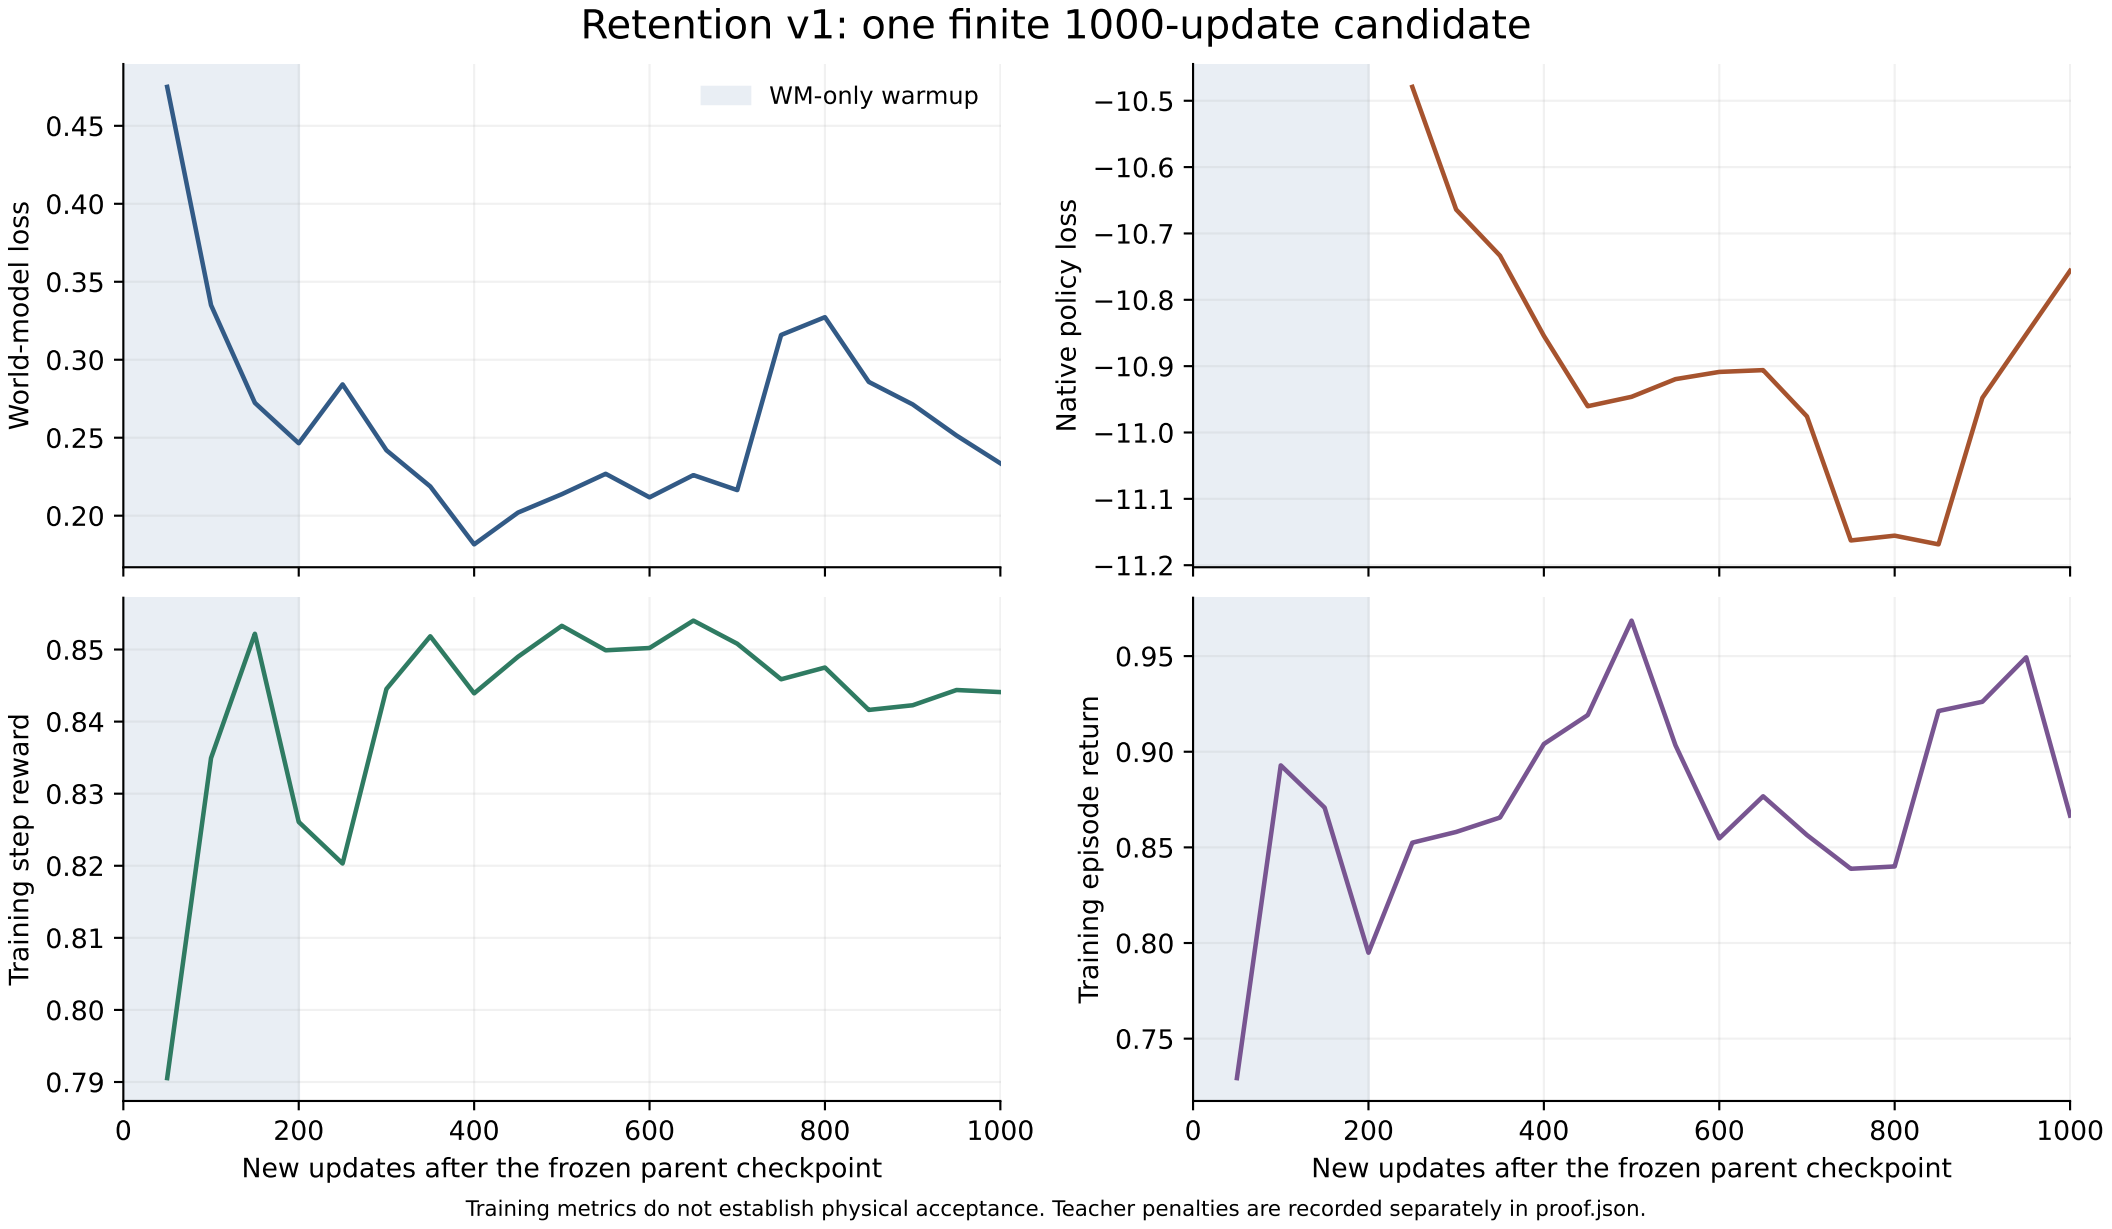
Retention v1: one finite 1000-update candidate
WM-only warmup
World-model loss
Native policy loss
Training step reward
Training episode return
New updates after the frozen parent checkpoint
New updates after the frozen parent checkpoint
Training metrics do not establish physical acceptance. Teacher penalties are recorded separately in proof.json.